Discovering activation functions between layers

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Maham Faisal Khan
Senior Data Science Content Developer

datacamp

Limitations of the sigmoid and softmax function

Sigmoid functions:

- Bounded between 0 and 1
- Can be used anywhere in the network

Gradients:

- Approach zero for low and high values of x
- Cause function to **saturate**

Sigmoid function saturation can lead to **vanishing gradients** during backpropagation.

This is also a problem for **softmax**.

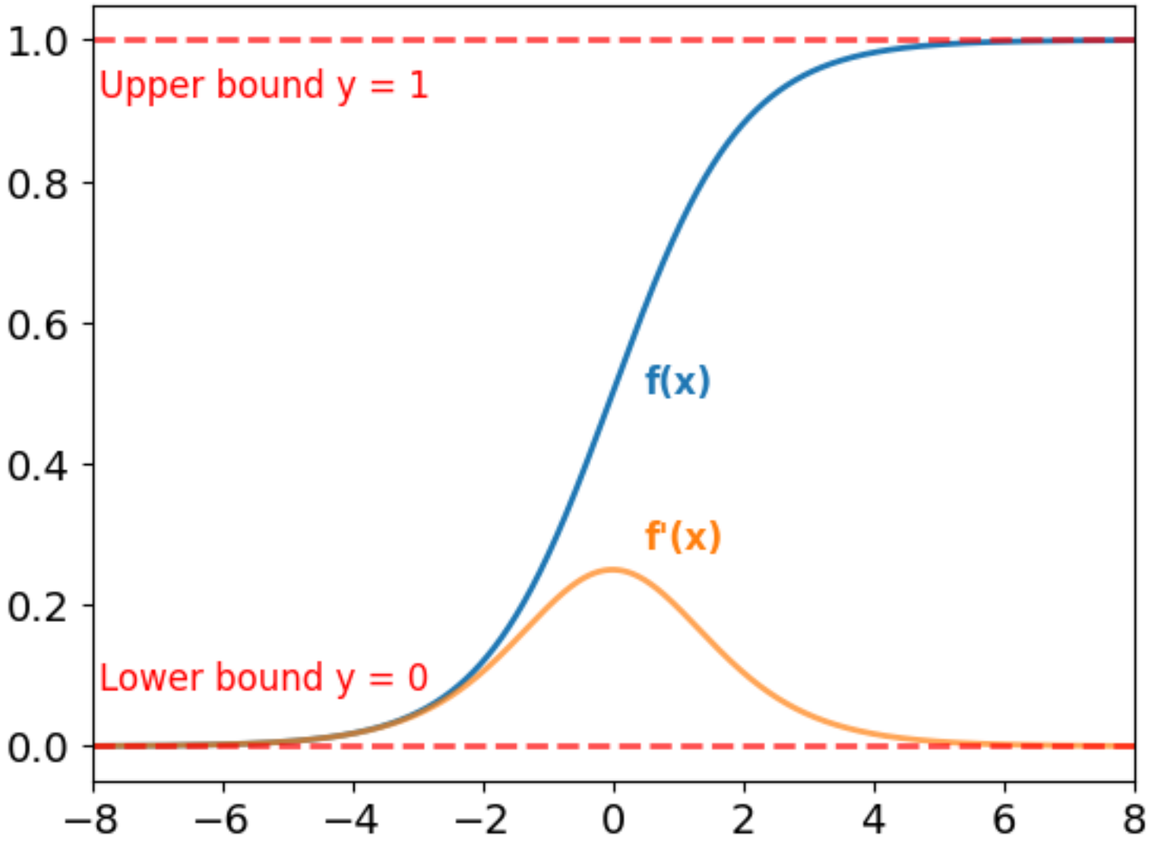

Introducing ReLU

Rectified Linear Unit (ReLU):

- `f(x) = max(x, 0)`
- for positive inputs, the output is equal to the input
- for strictly negative inputs, the output is equal to zero
- overcomes the vanishing gradients problem

In PyTorch:

```
relu = nn.ReLU()
```

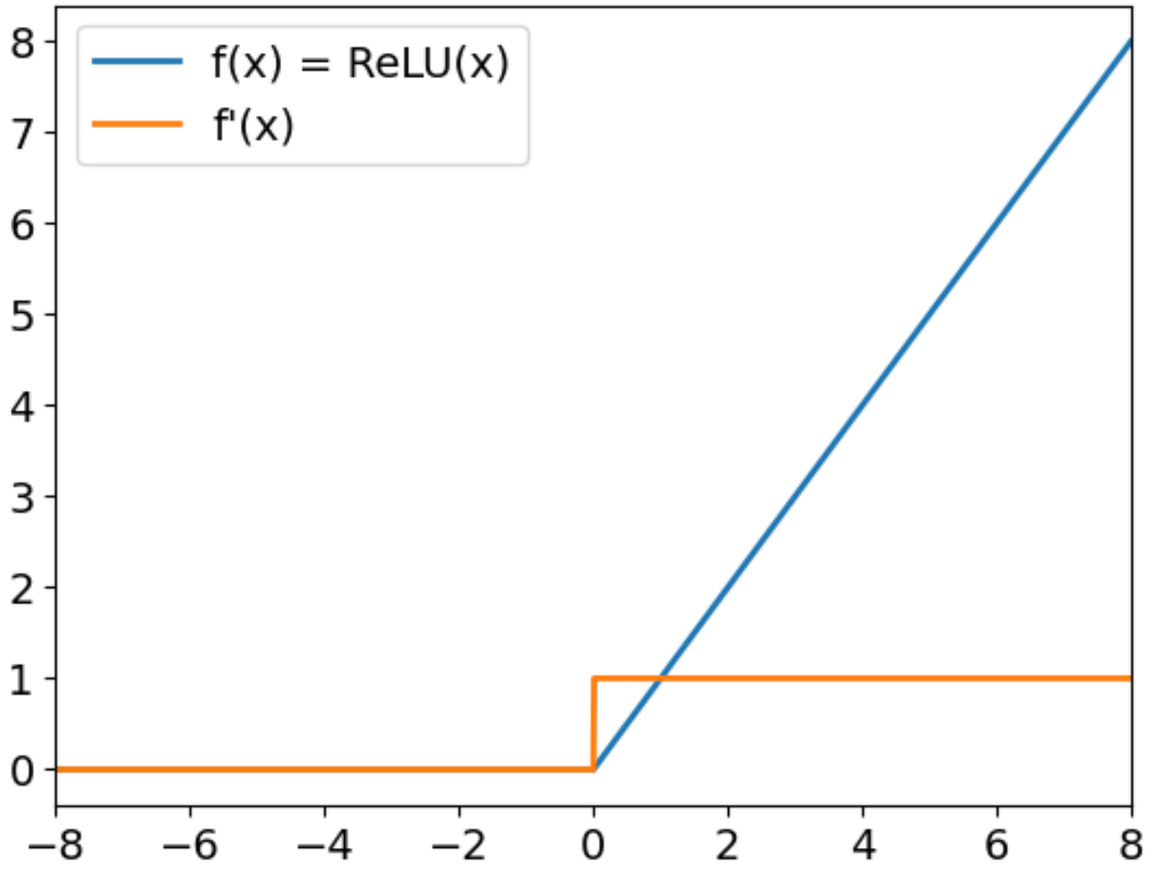

Introducing Leaky ReLU

Leaky ReLU:

- For positive inputs, it behaves similarly to ReLU
- For negative inputs, it multiplies the input by a small coefficient (defaulted to 0.01)
- The gradients for negative inputs are never null

In PyTorch:

```
leaky_relu = nn.LeakyReLU(negative_slope = 0.05)
```

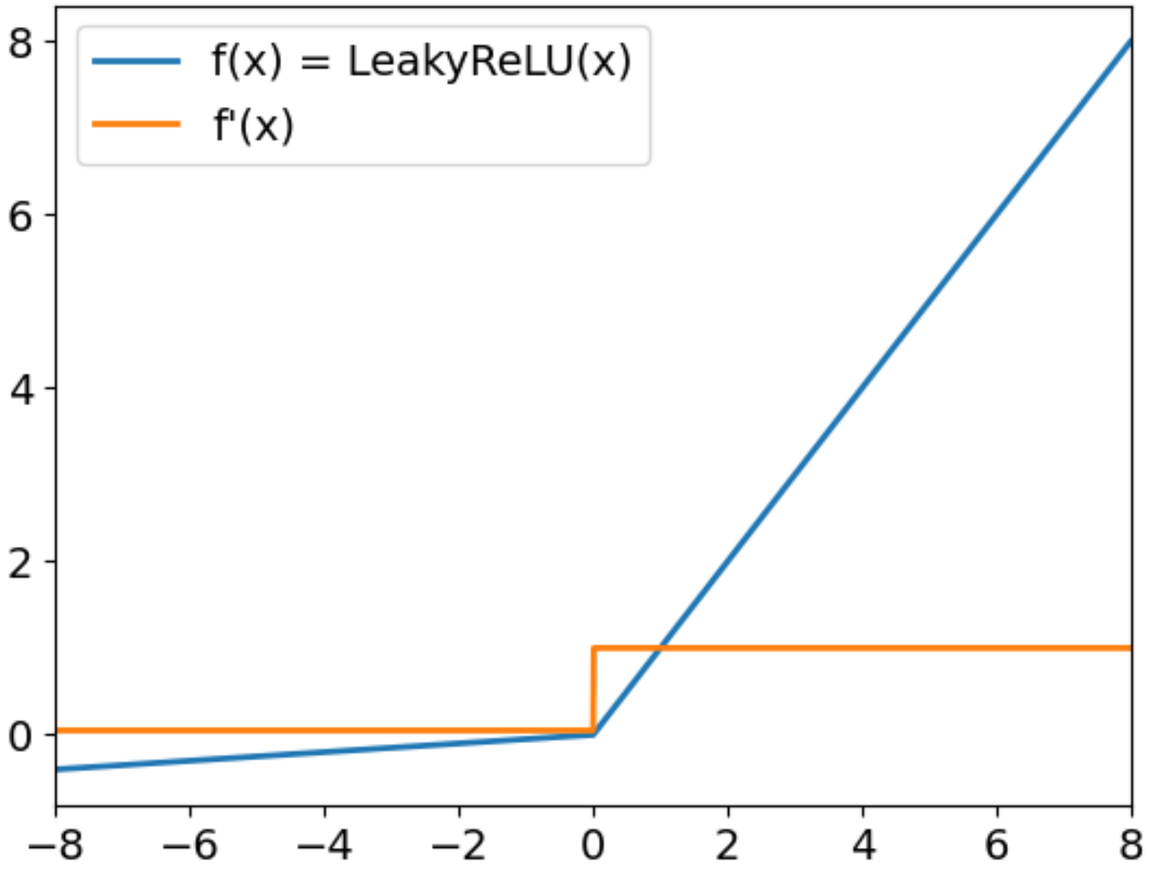

Let's practice!

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

A deeper dive into neural network architecture

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Maham Faisal Khan
Senior Data Science Content Developer

datacamp

Layers are made of neurons

- Linear layers are **fully connected**
- Each neuron of a layer connected to each neuron of previous layer

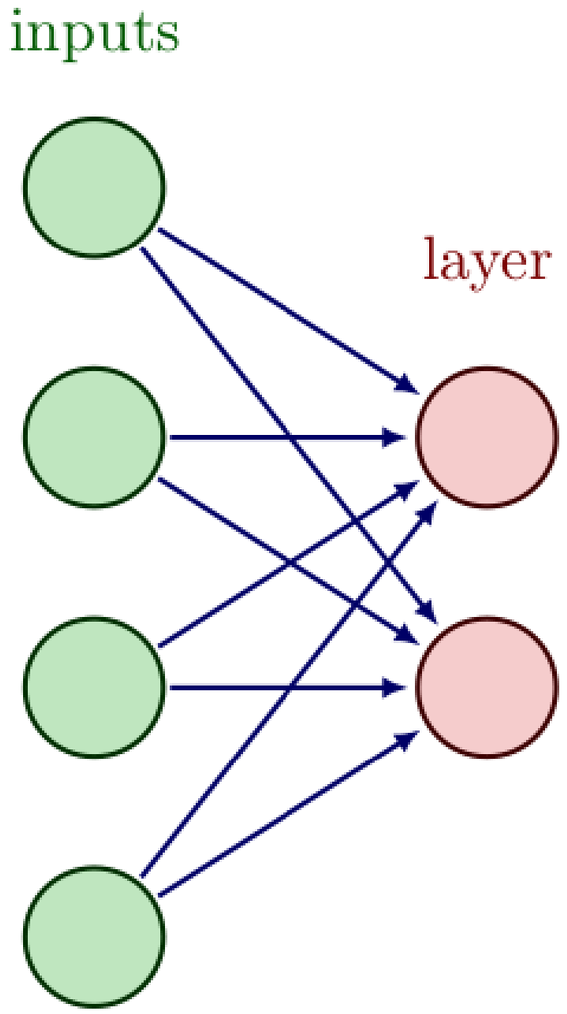

- A neuron of a linear layer:
  - computes a linear operation using all neurons of previous layer
  - contains N+1 learnable parameters
  - where N = dimension of previous layer's outputs

Layer naming convention

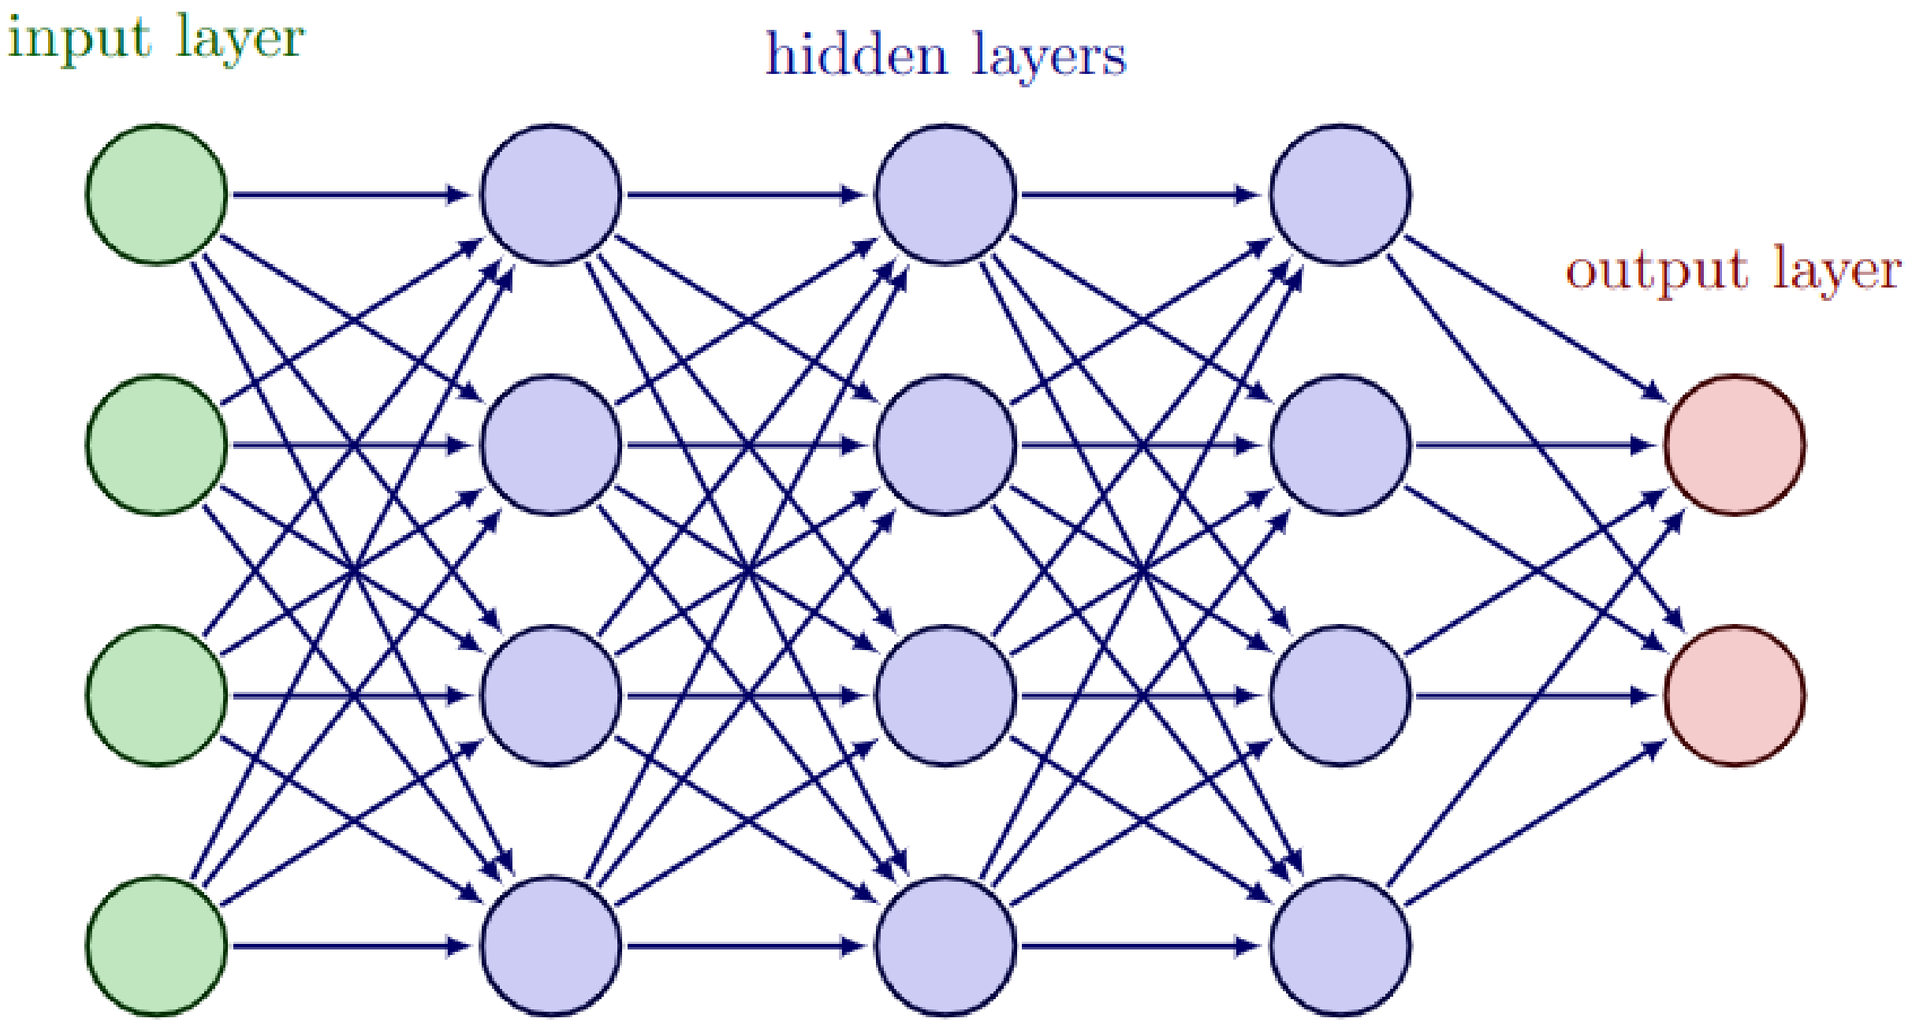

Tweaking the number of hidden layers

- Input and output layers dimensions are fixed.
  - input layer depends on the number of features `n_features`
  - output layer depends on the number of categories `n_classes`

```
model = nn.Sequential(nn.Linear(n_features, 8),
                      nn.Linear(8, 4),
                      nn.Linear(4, n_classes))
```

- We can use as many hidden layers as we want
- Increasing the number of hidden layers = increasing the number of parameters = increasing the **model capacity**

Counting the number of parameters

Given the following model:

```
model = nn.Sequential(nn.Linear(8, 4),
                      nn.Linear(4, 2))
```

Manually calculating the number of parameters:

- first layer has 4 neurons, each neuron has 8+1 parameters = 36 parameters
- second layer has 2 neurons, each neuron has 4+1 parameters = 10 parameters
- total = 46 learnable parameters

Using PyTorch:

- `.numel()` : returns the number of elements in the tensor

```
total = 0
for parameter in model.parameters():
    total += parameter.numel()
print(total)
```

```
46
```

Let's practice!

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Learning rate and momentum

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Maham Faisal Khan
Senior Data Science Content Developer

datacamp

Updating weights with SGD

- Training a neural network = solving an **optimization problem.**

Stochastic Gradient Descent (SGD) optimizer

```
sgd = optim.SGD(model.parameters(), lr=0.01, momentum=0.95)
```

- Two parameters:
  - **learning rate:** controls the step size
  - **momentum:** controls the inertia of the optimizer
- Bad values can lead to:
  - long training times
  - bad overall performances (poor accuracy)

Impact of the learning rate: optimal learning rate

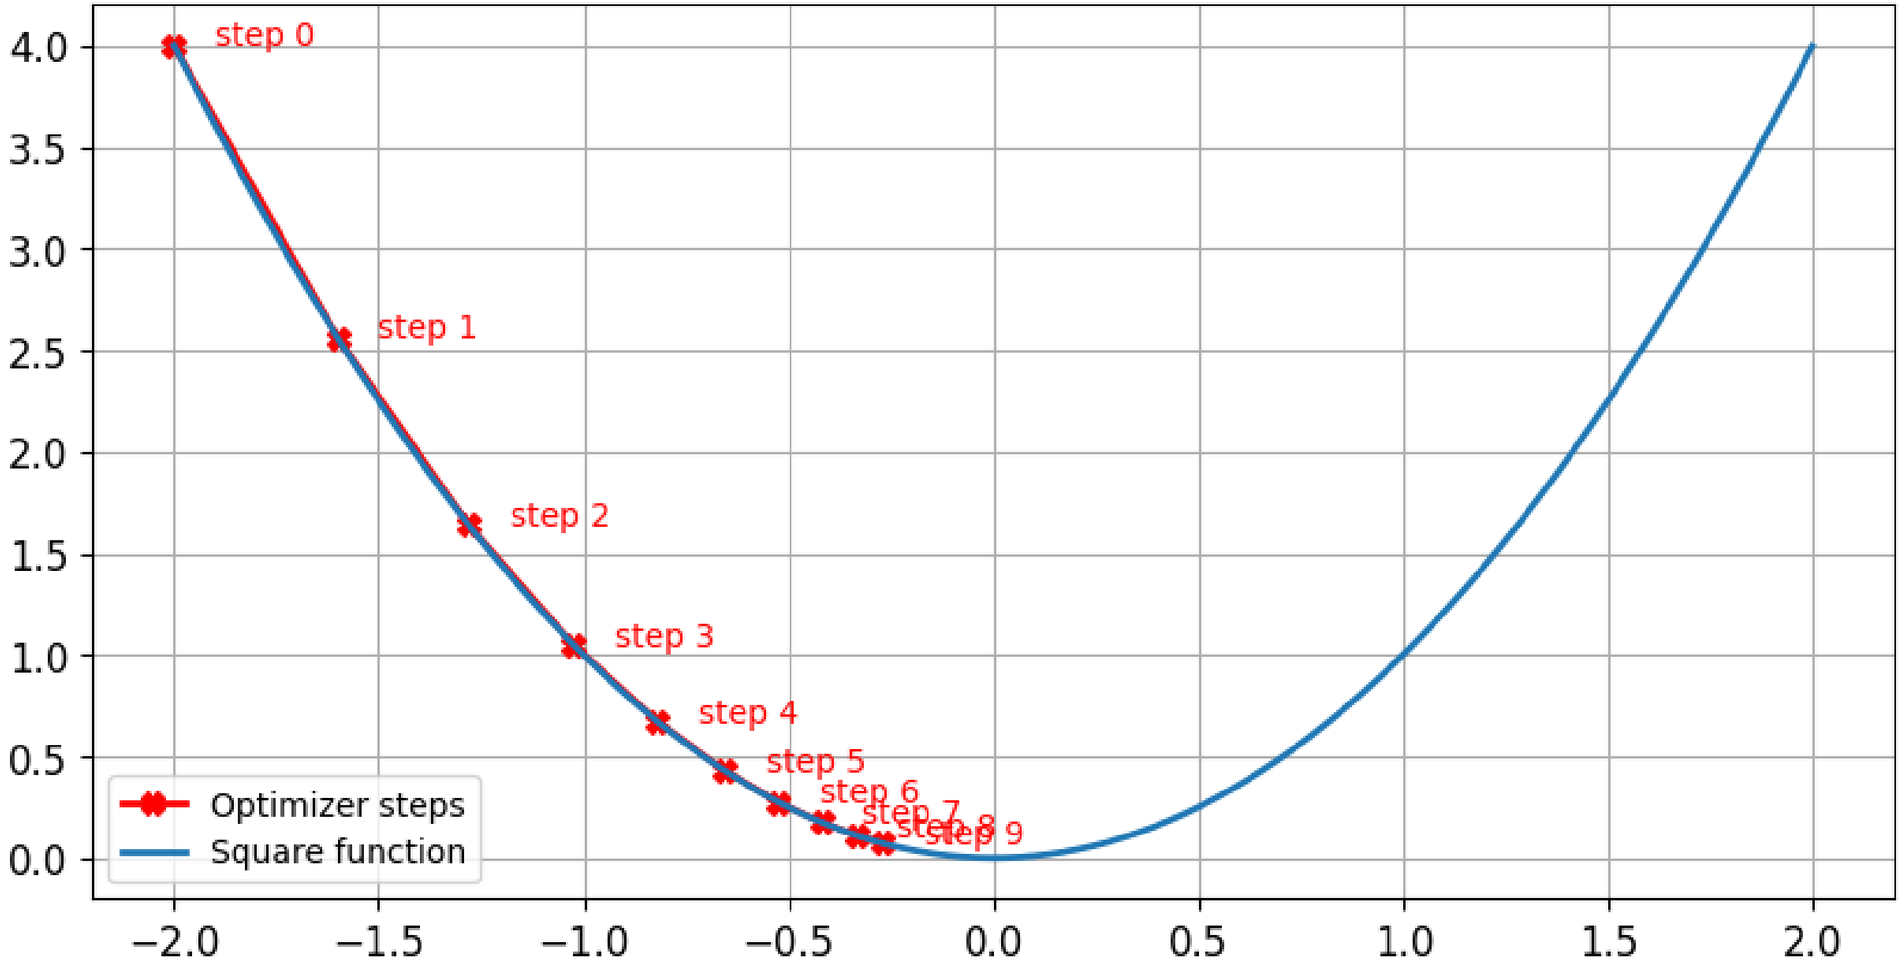

Impact of the learning rate: small learning rate

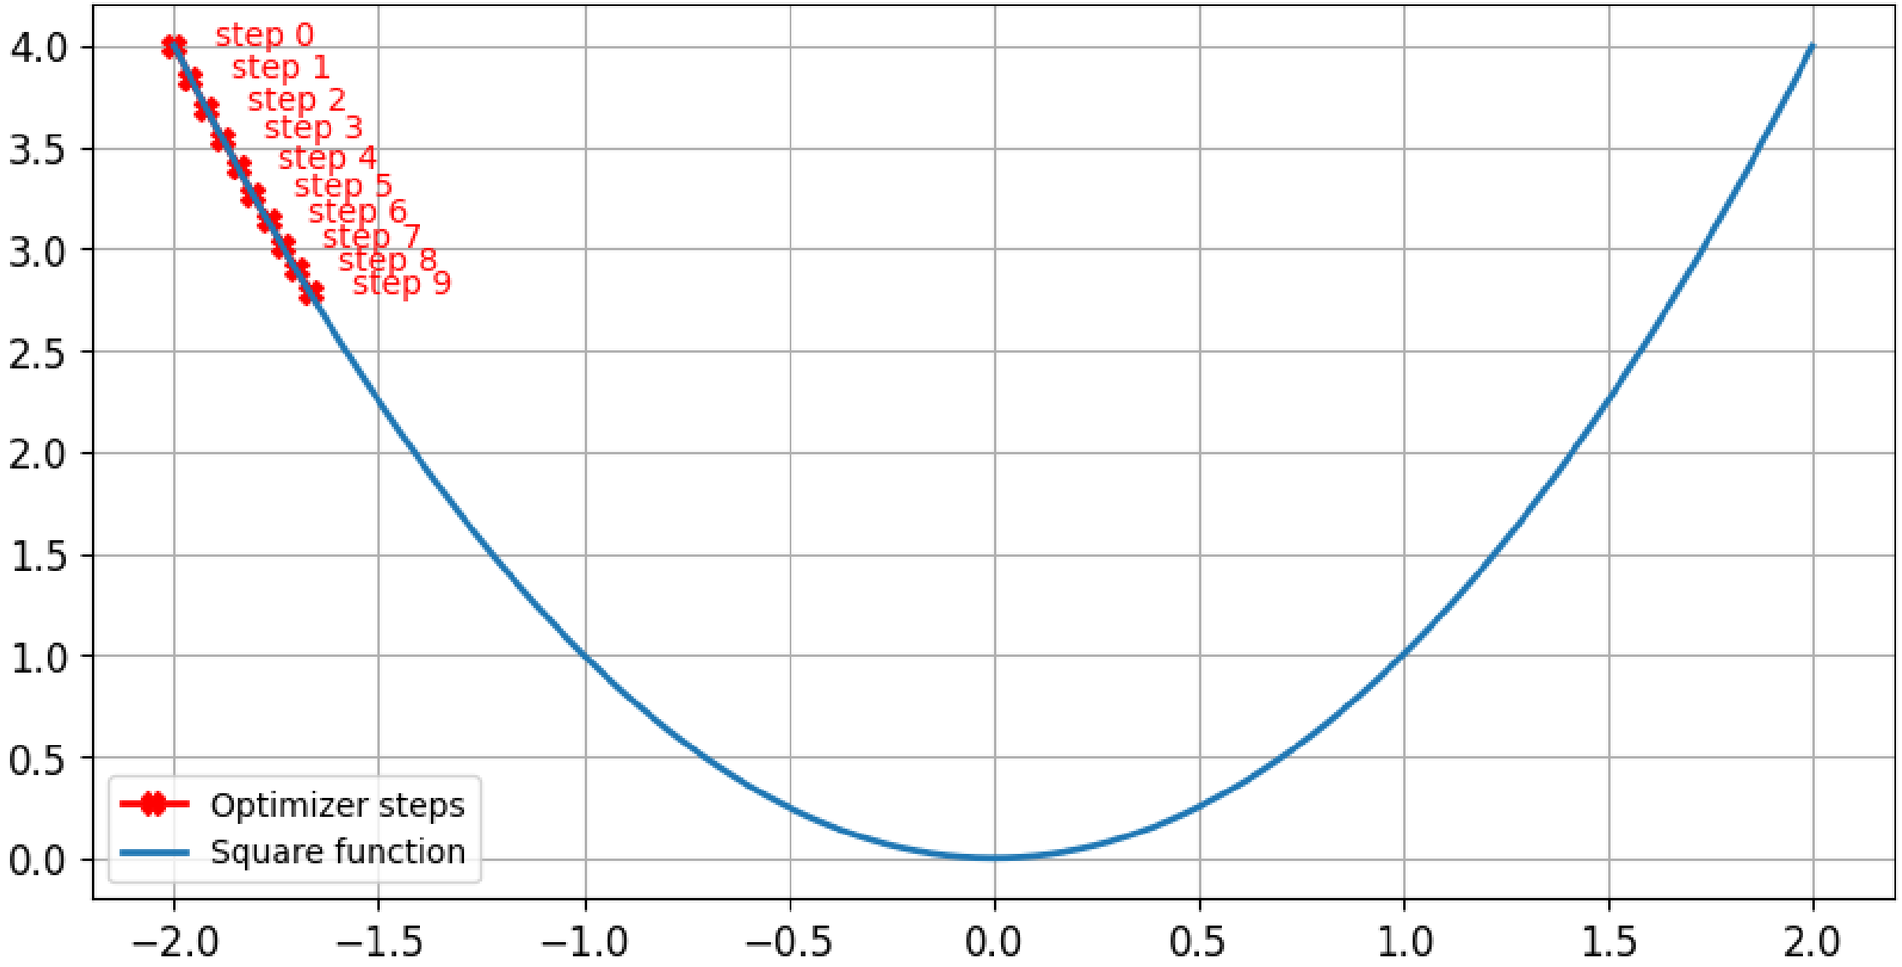

Impact of the learning rate: high learning rate

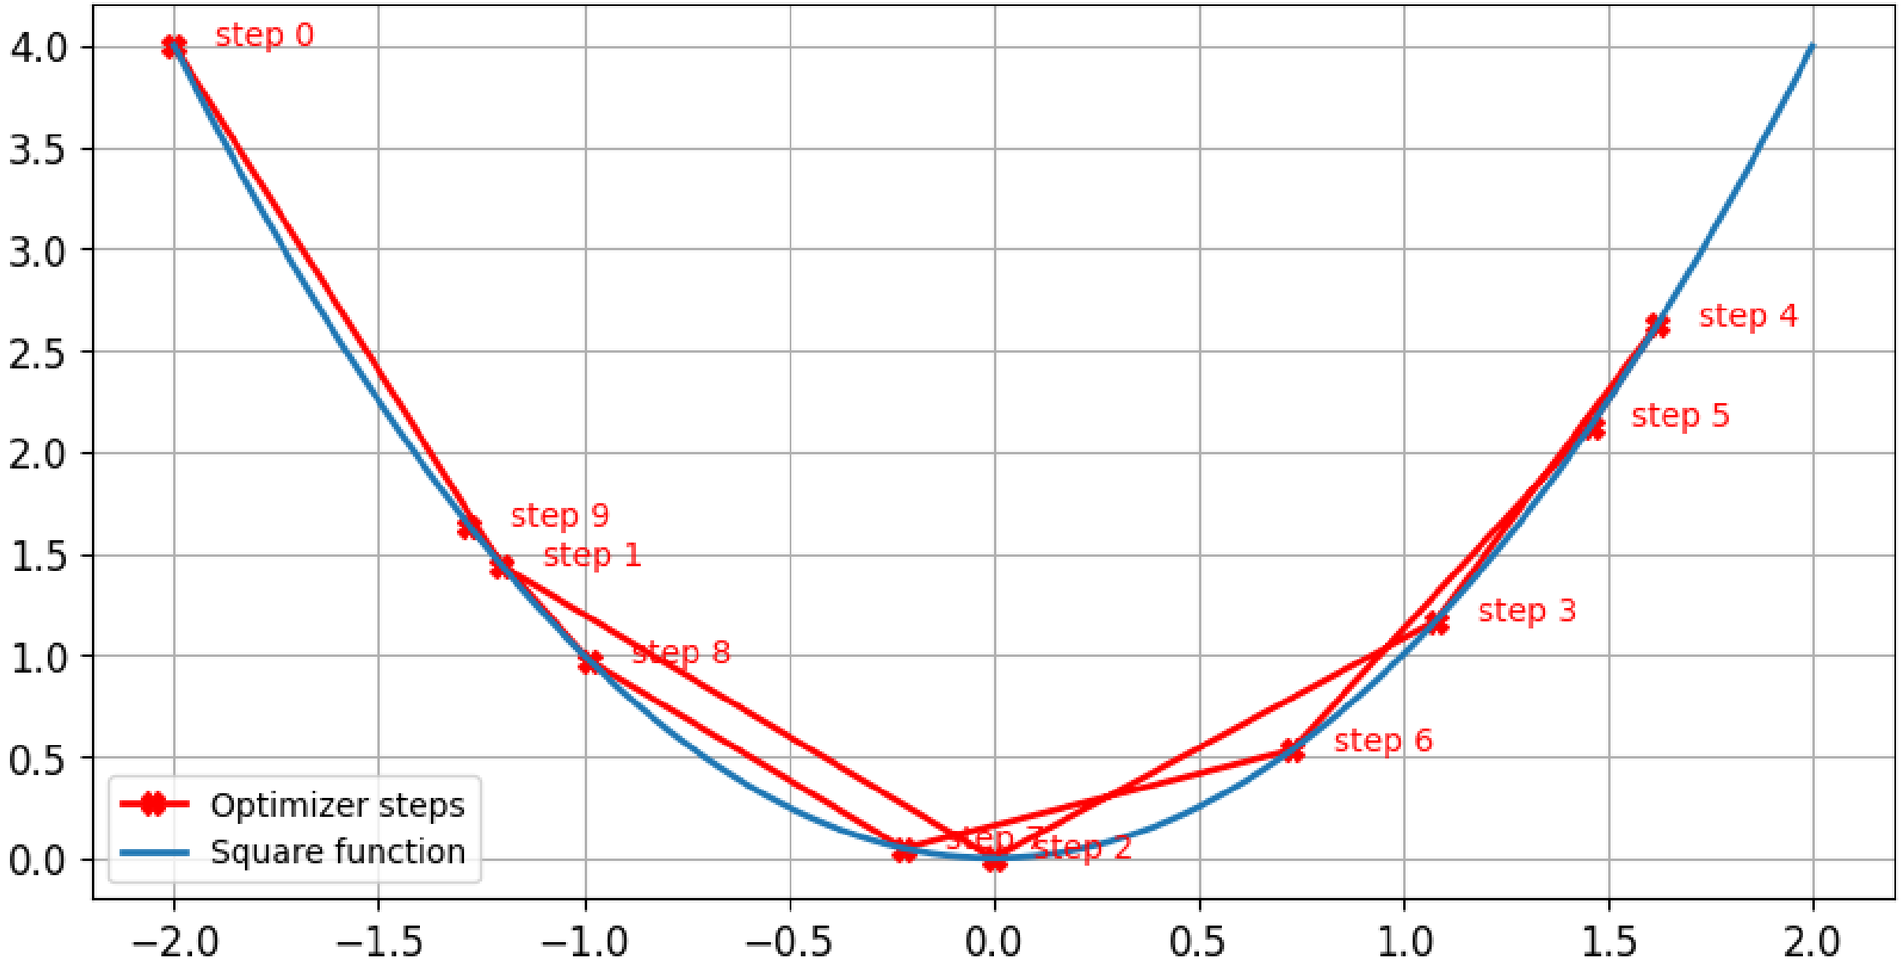

Without momentum

- `lr = 0.01` `momentum = 0` , after 100 steps minimum found for `x = -1.23` and `y = -0.14`

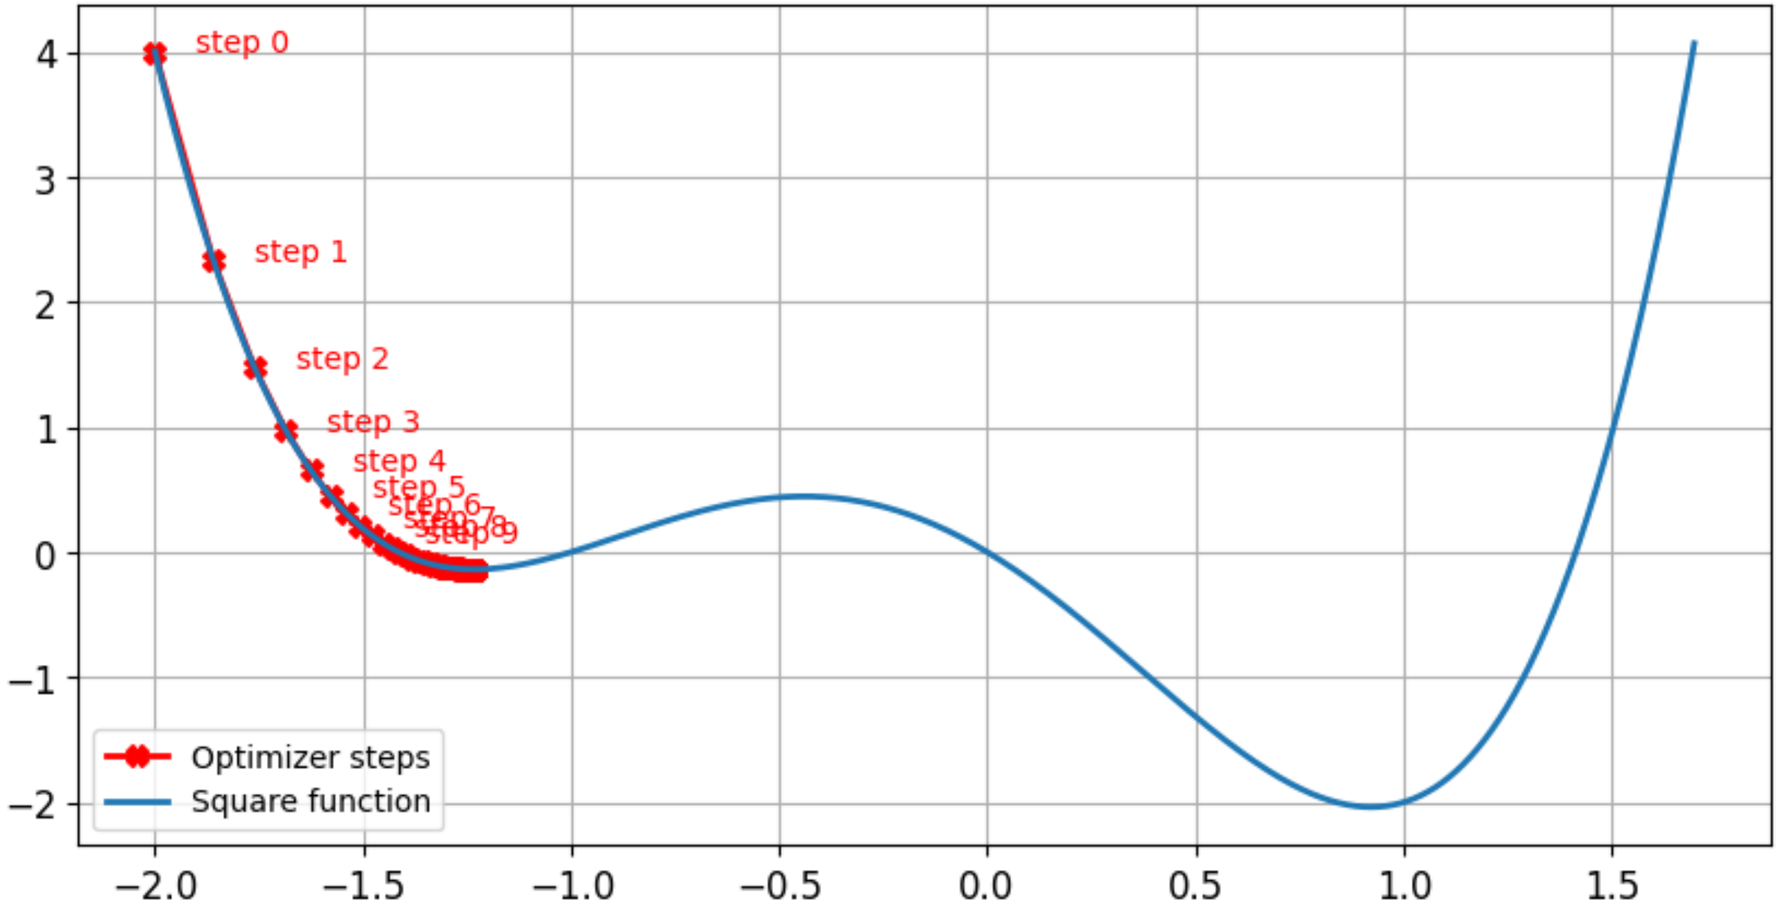

With momentum

- `lr = 0.01` `momentum = 0.9` , after 100 steps minimum found for `x = 0.92` and `y = -2.04`

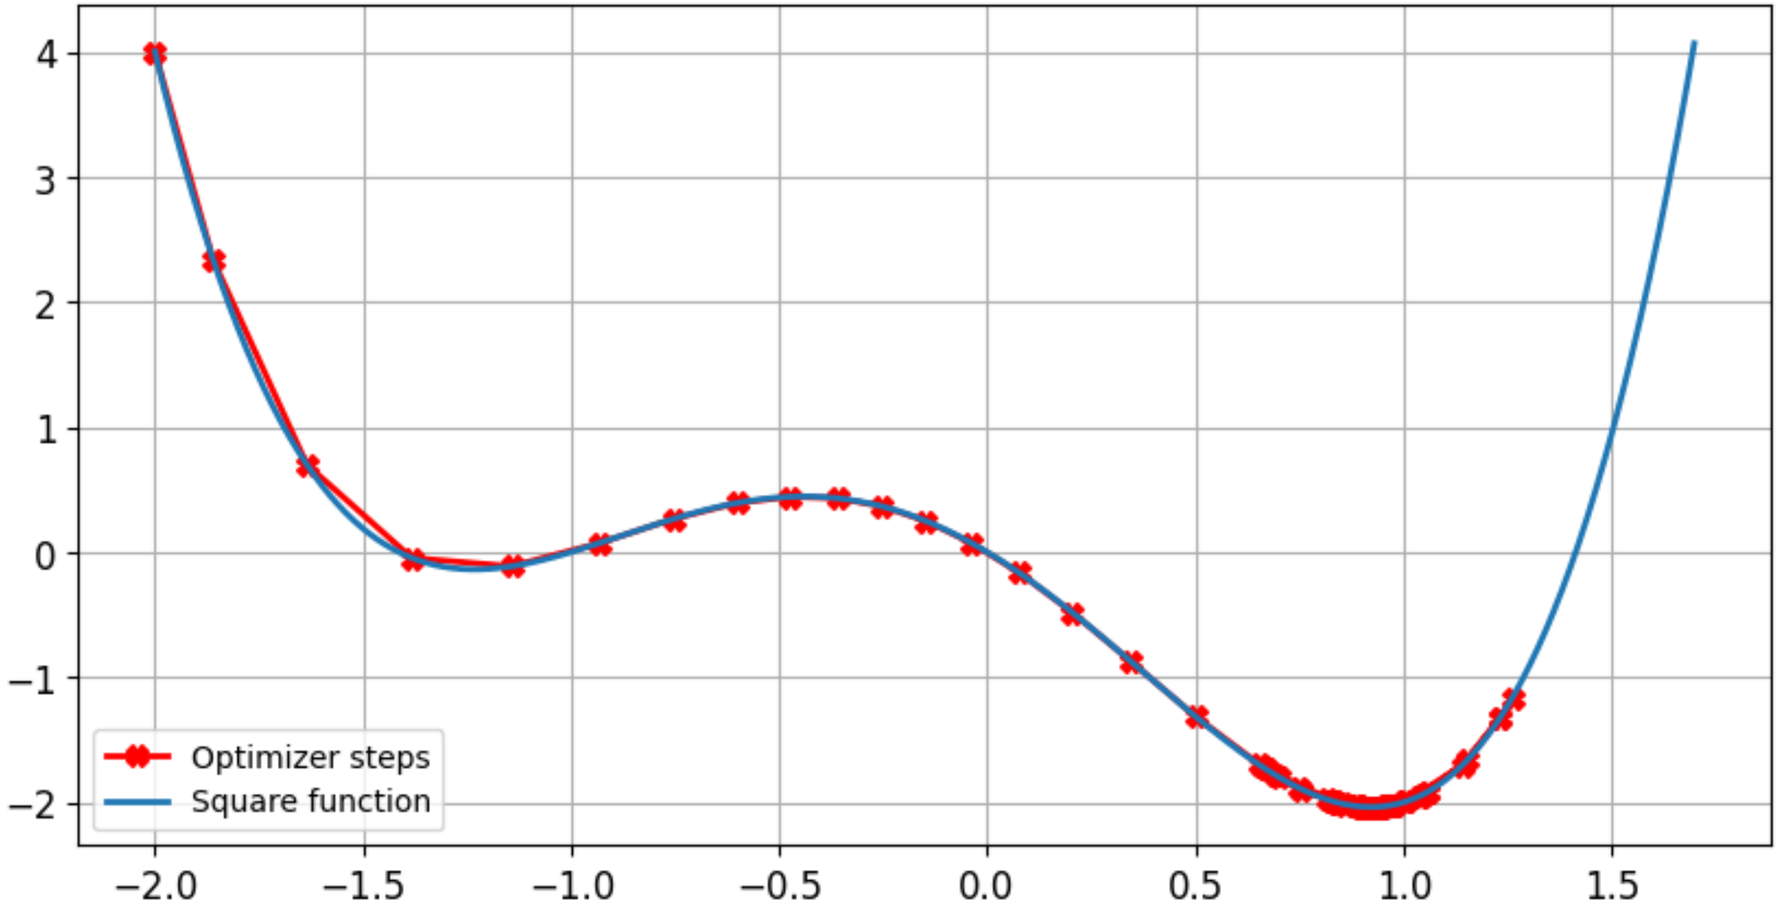

Summary

Learning rate	Momentum
Controls the step size	Controls the inertia
Too small leads to long training times	Null momentum can lead to the optimizer being stuck in a local minimum
Too high leads to poor performances	Non-null momentum can help find the function minimum
Typical values between 10^{-2} and 10^{-4}	Typical values between 0.85 and 0.99

Let's practice!

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Layer initialization and transfer learning

INTRODUCTION TO DEEP LEARNING WITH PYTORCH

Maham Faisal Khan
Senior Data Science Content Developer

datacamp

Layer initialization

```
import torch.nn as nn
layer = nn.Linear(64, 128)
print(layer.weight.min(), layer.weight.max())
```

```
(tensor(-0.1250, grad_fn=<MinBackward1>), tensor(0.1250, grad_fn=<MaxBackward1>))
```

- Layer weights are initialized to small values
- Layer outputs can explode if inputs and weights are not normalized
- Weights can be initialized using different methods (e.g., with a uniform distribution)

Layer initialization

```
import torch.nn as nn


layer = nn.Linear(64, 128)
nn.init.uniform_(layer.weight)
print(custom_layer.fc.weight.min(), custom_layer.fc.weight.max())
```

```
(tensor(0.0002, grad_fn=<MinBackward1>), tensor(1.0000, grad_fn=<MaxBackward1>))
```

Transfer learning and fine-tuning

Transfer learning: reusing a model trained on a first task for a second similar task, to accelerate the training process.

```
import torch

layer = nn.Linear(64, 128)
torch.save(layer, 'layer.pth')

new_layer = torch.load('layer.pth')
```

Transfer learning and fine-tuning

- **Fine-tuning** = A type of transfer learning
  - Smaller learning rate
  - Not every layer is trained (we **freeze** some of them)
  - Rule of thumb: freeze early layers of network and fine-tune layers closer to output layer

```
import torch.nn as nn

model = nn.Sequential(nn.Linear(64, 128),
                      nn.Linear(128, 256))

for name, param in model.named_parameters():
    if name == '0.weight':
        param.requires_grad = False
```

Let's practice!

INTRODUCTION TO DEEP LEARNING WITH PYTORCH